

Data Mining Assignment 5

Xuan Han
han.xua@husky.neu

October 12, 2015

8: Cross validation


(a)

```
> set.seed(1)
> y = rnorm(100)
> x = rnorm(100)
> y = x - 2 * x^2 + rnorm(100)
```

1. In this data set, $n = 100$, $p = 2$
2. $Y = X - 2X^2 + \epsilon$

(b)

```
> plot(y ~ x, col = 'blue')
```



1. We can see that the data points are distributed along a parabola, just as we expected.
2. x goes from -2 to 2, y goes from -8 to 2.

(c)

```
> library(boot)
> set.seed(12)
> df = data.frame(X = x, Y = y)
> p = 4
> loocv.err = rep(NA, p)
> for(i in 1:p) {
+   glm.fit = glm(Y ~ poly(X, i), data = df)
+   loocv.err[i] = cv.glm(df, glm.fit)$delta[1]
+ }
> loocv.err
```

```
[1] 5.890979 1.086596 1.102585 1.114772
```

(d)

```
> set.seed(2015)
> loocv.err = rep(NA, p)
> for(i in 1:p) {
+   glm.fit = glm(Y ~ poly(X, i), data = df)
+   loocv.err[i] = cv.glm(df, glm.fit)$delta[1]
+ }
> loocv.err
```

```
[1] 5.890979 1.086596 1.102585 1.114772
```

1. The results obtained are exactly the same.
2. The reason is that for leave-one-out-CV, the results have nothing to do with seed. It will evaluate 100 times no matter what the seed is. There is no randomness involved in the process of LOOCV.

(e)

1. The second model have the smallest LSSCV error. This is exactly what I expected.
2. Because the simulated data set is generated by power of 2 of X , if we use a lower order model such as power of 1 of X , it will have high bias; If we use a higher order model such as power of 3 or 4 of X , it will cause high variance. So, power of 2 have low bias and low variance.

(f)

```
> glm.fit = glm(Y ~ poly(X, 1), data = df)
> summary(glm.fit)$coeff
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.827707	0.2362206	-7.7372898	9.181461e-12
poly(X, 1)	2.316401	2.3622062	0.9806091	3.292002e-01

```
> glm.fit = glm(Y ~ poly(X, 2), data = df)
> summary(glm.fit)$coeff
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.827707	0.1032351	-17.70431	3.804657e-32
poly(X, 2)1	2.316401	1.0323515	2.24381	2.711854e-02
poly(X, 2)2	-21.058587	1.0323515	-20.39866	7.333860e-37

```
> glm.fit = glm(Y ~ poly(X, 3), data = df)
> summary(glm.fit)$coeff
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.8277074	0.1037248	-17.6207390	7.610579e-32
poly(X, 3)1	2.3164010	1.0372479	2.2332183	2.785714e-02
poly(X, 3)2	-21.0585869	1.0372479	-20.3023667	1.636959e-36
poly(X, 3)3	-0.3048398	1.0372479	-0.2938929	7.694742e-01

```
> glm.fit = glm(Y ~ poly(X, 4), data = df)
> summary(glm.fit)$coeff
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.8277074	0.1041467	-17.5493533	1.444977e-31
poly(X, 4)1	2.3164010	1.0414671	2.2241711	2.850549e-02
poly(X, 4)2	-21.0585869	1.0414671	-20.2201171	3.457023e-36
poly(X, 4)3	-0.3048398	1.0414671	-0.2927023	7.703881e-01
poly(X, 4)4	-0.4926249	1.0414671	-0.4730105	6.372907e-01

```
>
```

1. From the coefficient estimates we can see that liner and quradratic term of X have vary low p-value while cubic and higher order have high p-value. This results agree with the conclusion dran based on the CV results.

9: Surgical

(a)

```
> X <- read.table('surgical.txt')
> dimnames(X)[[2]] <- c('blood', 'prog', 'enz', 'liver', 'age',
+                        'female', 'modAlc', 'heavyAlc', 'surv', 'lsurv')
> X = X[, -9]
> N = dim(X)[1]
> d = dim(X)[2]
> mu.hat = mean(X$lsurv)
> mu.hat
```

```
[1] 6.430481
```

(b)

```
> se = sd(X$lsurv) / sqrt(N)
> se
```

```
[1] 0.06689616
```

1. This result suggests that the expected deviation of sample mean of lsurv is 0.875

(c)

```
> boot.fn = function(data, index) return (mean(data[index]))
> boot(X$lsurv, boot.fn, R = 1000)
```

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:

```
boot(data = X$lsurv, statistic = boot.fn, R = 1000)
```

Bootstrap Statistics :

	original	bias	std. error
t1*	6.430481	-0.003804222	0.06627289

1. We can see that the results are very close. Before se is 0.066896, now it is 0.066885

(d)

```
> confd.low = mu.hat - 2 * se
> confd.high = mu.hat + 2 * se
> c(confd.low, confd.high)
```

```
[1] 6.296689 6.564274
```

```
> t.test(X$lsurv)
```

One Sample t-test

```
data: X$lsurv
t = 96.126, df = 53, p-value < 2.2e-16
alternative hypothesis: true mean is not equal to 0
```

```

95 percent confidence interval:
 6.296305 6.564658
sample estimates:
mean of x
 6.430481

```

1. It is clear that the results are very close.

(e)

```

> md.hat = median(X$lurv)
> md.hat

```

```

[1] 6.406

```

(f)

```

> boot.fn = function(data, index) return(median(data[index]))
> boot(X$lurv, boot.fn, R = 1000)

```

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:

```
boot(data = X$lurv, statistic = boot.fn, R = 1000)
```

Bootstrap Statistics :

	original	bias	std. error
t1*	6.406	0.009343	0.05701799

1. Median estimated by bootstrap is 6.046 which is exactly the same above.
2. The std.error estimated by bootstrap is 0.0576, which is small std error.

(g)

```

> quantile(X$lurv, 0.1)

```

```

10%
5.8711

```

(h)

```

> boot.fn = function(data, index) return(quantile(data[index], 0.1))
> boot(X$lurv, boot.fn, R = 1000)

```

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:

```
boot(data = X$lurv, statistic = boot.fn, R = 1000)
```

Bootstrap Statistics :

	original	bias	std. error
t1*	5.8711	0.011588	0.113706

1. The tenth quantile is the same with above.
2. Here we get a small std error for the tenth quantile, which is 0.11

9: Ridge and Lasso

(a)

```
> set.seed(1)
> test = sample(1:N, N / 2)
> train = -test
> X.train = X[train,]
> X.test = X[test, ]
```

(b)

```
> lm.fit = lm(lsurv ~ ., data = X.train)
> mrs = mean((predict.lm(lm.fit, X.test) - X.test$lsurv) ^ 2)
> mrs
```

```
[1] 0.05299499
```


(c)

```
> library(glmnet)
> x.train = model.matrix(lsurv ~ ., X.train)[, -1]
> y.train = X.train$lsurv
> grid = 10 ^ seq(10, -2, length = 100)
> ridge.mod = glmnet(x.train, y.train, alpha = 0, lambda = grid, thresh = 1e-12)
> cv.out = cv.glmnet(x.train, y.train, alpha = 0, nfolds = 10)
> plot(cv.out)
> bestlam = cv.out$lambda.min
> bestlam
```

```
[1] 0.06623395
```

```
> x.test = model.matrix(lsurv ~ ., X.test)[, -1]
> y.test = X.test$lsurv
> ridge.pred = predict(ridge.mod, s = bestlam, newx = x.test)
> ridge.mrs = mean((ridge.pred - y.test)^2)
> ridge.mrs
```

```
[1] 0.06356988
```



(d)

```
> ridge.mod = glmnet(x.train, y.train, alpha = 1, lambda = grid, thresh = 1e-12)
> cv.out = cv.glmnet(x.train, y.train, alpha = 1)
> plot(cv.out)
> bestlam = cv.out$lambda.min
> bestlam
```

```
[1] 0.008357777
```


```
> ridge.pred = predict(ridge.mod, s = bestlam, newx = x.test)
> lasso.mrs = mean((ridge.pred - y.test) ^ 2)
> lasso.mrs
```

```
[1] 0.05464687
```

```
> lasso.coef = predict(cv.out, type = 'coeff', s = bestlam)
> lasso.coef
```

```
9 x 1 sparse Matrix of class "dgCMatrix"
```

```
      1
(Intercept) 4.370137313
blood      0.050607955
prog       0.009716315
enz        0.011266202
liver      0.116956233
age        -0.002251710
female     0.024308981
modAlc     .
heavyAlc   0.391001280
```



(g)

1. We get a test error of 0.053 for linear model; 0.064 for ridge; 0.054 for lasso
2. The test error get from ridge and lasso are higher than test error directly from a linear mode.
3. From lass we see that coeff of predictor modAlc is set zero.

```
> y.test.mu = mean(y.test)
> lm.r2 = 1 - mrs / mean((y.test - y.test.mu) ^ 2)
> ridge.r2 = 1 - ridge.mrs / mean((y.test - y.test.mu) ^ 2)
> lasso.r2 = 1 - lasso.mrs / mean((y.test - y.test.mu) ^ 2)
> lm.r2
```

```
[1] 0.7983984
```

```
> ridge.r2
```

```
[1] 0.7581698
```

```
> lasso.r2
```

```
[1] 0.7921144
```

1. From the R^2 statistic we can see that these three model only explained 79% of the variance of the data, so the accuracy of our model is moderate.